

Python For Data Science Cheat Sheet

PySpark Basics

Learn Python for data science interactively at [www.DataCamp.com](https://www.datacamp.com)



Spark

PySpark is the Spark Python API that exposes the Spark programming model to Python



Initializing Spark

SparkContext

```
>>> from pyspark import SparkContext
>>> sc = SparkContext(master = 'local[2]')
```

Inspect SparkContext

>>> sc.version	Retrieve SparkContext version
>>> sc.pythonVer	Retrieve Python version
>>> sc.master	Master URL to connect to
>>> str(sc.sparkHome)	Path where Spark is installed on worker nodes
>>> str(sc.sparkUser())	Retrieve name of the Spark User running SparkContext
>>> sc.appName	Return application name
>>> sc.applicationId	Retrieve application ID
>>> sc.defaultParallelism	Return default level of parallelism
>>> sc.defaultMinPartitions	Default minimum number of partitions for RDDs

Configuration

```
>>> from pyspark import SparkConf, SparkContext
>>> conf = (SparkConf()
>>> .setMaster("local")
>>> .setAppName("My app")
>>> .set("spark.executor.memory", "1g"))
>>> sc = SparkContext(conf = conf)
```

Using The Shell

In the PySpark shell, a special interpreter-aware SparkContext is already created in the variable called `sc`.

```
$ ./bin/spark-shell --master local[2]
$ ./bin/pyspark --master local[4] --py-files code.py
```

Set which master the context connects to with the `--master` argument, and add Python `.zip`, `.egg` or `.py` files to the runtime path by passing a comma-separated list to `--py-files`.

Loading Data

Parallelized Collections

```
>>> rdd = sc.parallelize([('a',7), ('a',2), ('b',2)])
>>> rdd2 = sc.parallelize([('a',2), ('d',1), ('b',1)])
>>> rdd3 = sc.parallelize(range(100))
>>> rdd4 = sc.parallelize(["a","x","y","z"],
>>>                        ("b","p","r"]])
```

External Data

Read either one text file from HDFS, a local file system or any Hadoop-supported file system URI with `textFile()`, or read in a directory of text files with `wholeTextFiles()`.

```
>>> textFile = sc.textFile("/my/directory/*.txt")
>>> textFile2 = sc.wholeTextFiles("/my/directory/")
```

Retrieving RDD Information

Basic Information

>>> rdd.getNumPartitions()	List the number of partitions
>>> rdd.count()	Count RDD instances
>>> rdd.countByKey()	Count RDD instances by key
>>> rdd.countByValue()	Count RDD instances by value
>>> rdd.collectAsMap()	Return (key,value) pairs as a dictionary
>>> rdd3.sum()	Sum of RDD elements
>>> sc.parallelize([]).isEmpty()	Check whether RDD is empty

Summary

>>> rdd3.max()	Maximum value of RDD elements
>>> rdd3.min()	Minimum value of RDD elements
>>> rdd3.mean()	Mean value of RDD elements
>>> rdd3.stdev()	Standard deviation of RDD elements
>>> rdd3.variance()	Compute variance of RDD elements
>>> rdd3.histogram(3)	Compute histogram by bins
>>> rdd3.stats()	Summary statistics (count, mean, stdev, max & min)

Applying Functions

>>> rdd.map(lambda x: x*(x[1],x[0]))	Apply a function to each RDD element
>>> rdd5 = rdd.flatMap(lambda x: x*(x[1],x[0]))	Apply a function to each RDD element and flatten the result
>>> rdd5.collect()	Apply a flatMap function to each (key,value) pair of rdd4 without changing the keys

Selecting Data

>>> rdd.collect()	Return a list with all RDD elements
>>> rdd.take(2)	Take first 2 RDD elements
>>> rdd.first()	Take first RDD element
>>> rdd.top(2)	Take top 2 RDD elements
>>> rdd3.sample(False, 0.15, 81).collect()	Return sampled subset of rdd3
>>> rdd.filter(lambda x: "a" in x)	Filter the RDD
>>> rdd5.distinct().collect()	Return distinct RDD values
>>> rdd.keys().collect()	Return (key,value) RDD's keys

Iterating

>>> def g(x): print(x)	Apply a function to all RDD elements
>>> rdd.foreach(g)	
>>> rdd.foreach(g)	

Reshaping Data

>>> rdd.reduceByKey(lambda x,y: x+y)	Merge the rdd values for each key
>>> rdd.reduce(lambda a, b: a + b)	Merge the rdd values
>>> rdd3.groupBy(lambda x: x % 2)	Return RDD of grouped values
>>> rdd.groupByKey()	Group rdd by key
>>> seqOp = (lambda x,y: (x[0]+y,x[1]+1))	Aggregate RDD elements of each partition and then the results
>>> combOp = (lambda x,y: (x[0]+y[0],x[1]+y[1]))	Aggregate values of each RDD key
>>> rdd.aggregate((0,0),seqOp,combOp)	Aggregate the elements of each partition, and then the results
>>> rdd.aggregateByKey((0,0),seqOp,combOp)	Merge the values for each key
>>> rdd.foldByKey(0, add)	Create tuples of RDD elements by applying a function

Mathematical Operations

>>> rdd.subtract(rdd2)	Return each rdd value not contained in rdd2
>>> rdd2.subtractByKey(rdd)	Return each (key,value) pair of rdd2 with no matching key in rdd
>>> rdd.cartesian(rdd2).collect()	Return the Cartesian product of rdd and rdd2

Sort

>>> rdd2.sortBy(lambda x: x[1])	Sort RDD by given function
>>> rdd2.sortByKey()	Sort (key, value) RDD by key

Repartitioning

>>> rdd.repartition(4)	New RDD with 4 partitions
>>> rdd.coalesce(1)	Decrease the number of partitions in the RDD to 1

Saving

```
>>> rdd.saveAsTextFile("rdd.txt")
>>> rdd.saveAsHadoopFile("hdfs://namenodehost/parent/child",
>>>                      "org.apache.hadoop.mapred.TextOutputFormat")
```

Stopping SparkContext

```
>>> sc.stop()
```

Execution

```
$ ./bin/spark-submit examples/src/main/python/pi.py
```

DataCamp
Learn Python for Data Science interactively

